

The Pegasus Workflow Planner Properties

Property Documentation

V 2.0.1 automatically generated at

2007-07-19 15:32

Contents

1	Introduction	5
2	Property Files And Locations	6
2.1	pegasus.properties	6
2.2	pegasus.user.properties	6
3	Local Directories	7
3.1	pegasus.home.datadir	7
3.2	pegasus.home.sysconfdir	7
3.3	pegasus.home.sharedstatedir	7
3.4	pegasus.home.localstatedir	8
4	Site Directories	8
4.1	pegasus.dir.exec	8
4.2	pegasus.dir.storage	8
4.3	pegasus.dir.create	9
5	Schema File Location Properties	9
5.1	pegasus.schema.dax	9
5.2	pegasus.schema.pdax	10
5.3	pegasus.schema.sc	10
5.4	pegasus.schema.ivr	10

6	Database Drivers For All Relational Catalogs	11
6.1	pegasus.catalog.<catalog-name>.db.driver	11
6.2	pegasus.catalog.*.db.url	11
6.3	pegasus.catalog.*.db.user	12
6.4	pegasus.catalog.*.db.password	12
6.5	pegasus.catalog.*.db.*	12
6.6	pegasus.catalog.provenance	14
6.7	pegasus.catalog.work	15
7	Catalog Properties	15
7.1	Replica Catalog	15
7.1.1	pegasus.catalog.replica	15
7.1.2	pegasus.catalog.replica.url	16
7.1.3	pegasus.catalog.replica.chunk.size	17
7.1.4	pegasus.catalog.replica.lrc.ignore	17
7.1.5	pegasus.catalog.replica.lrc.restrict	17
7.1.6	pegasus.catalog.replica.cache.asrc	18
7.2	Site Catalog	18
7.2.1	pegasus.catalog.site	18
7.2.2	pegasus.catalog.site.file	19
7.3	Transformation Catalog	19
7.3.1	pegasus.catalog.transformation	19
7.3.2	pegasus.catalog.transformation.file	20
8	Replica Selection Properties	21
8.1	pegasus.selector.replica	21
8.2	pegasus.selector.replica.*.ignore.stagein.sites	22
8.3	pegasus.selector.replica.*.prefer.stagein.sites	22
9	Site Selection Properties	23
9.1	pegasus.selector.site	23
9.2	pegasus.selector.site.path	24
9.3	pegasus.site.selector.env.*	25
9.4	pegasus.selector.site.timeout	25
9.5	pegasus.selector.site.keep.tmp	25
10	Transfer Configuration Properties	26
10.1	pegasus.transfer.*.impl	26
10.2	pegasus.transfer.refiner	28
10.3	pegasus.transfer.arguments	29
10.4	pegasus.transfer.links	29
10.5	pegasus.transfer.force	30

10.6	pegasus.transfer.single.quote	30
10.7	pegasus.transfer.throttle.processes	30
10.8	pegasus.transfer.throttle.streams	31
10.9	pegasus.transfer.*.thirdparty.sites	31
10.10	pegasus.transfer.*.thirdparty.remote	32
10.11	pegasus.transfer.staging.delimiter	32
10.12	pegasus.transfer.disable.chmod.sites	33
10.13	pegasus.transfer.proxy	33
10.14	pegasus.transfer.*.priority	33
11	Gridstart And Exitcode Properties	34
11.1	pegasus.gridstart	34
11.2	pegasus.gridstart.invoke.always	34
11.3	pegasus.gridstart.invoke.length	35
11.4	pegasus.exitcode.scope	35
11.5	pegasus.exitcode.impl	36
11.6	pegasus.exitcode.path.[value]	37
11.7	pegasus.exitcode.arguments	38
11.8	pegasus.exitcode.debug	38
12	Remote Scheduler Properties	39
12.1	pegasus.remote.scheduler.projects	39
12.2	pegasus.remote.scheduler.queues	39
12.3	pegasus.remote.scheduler.min.[time]	39
13	Interface To Condor And Condor Dagman	40
13.1	pegasus.condor.arguments.quote	40
13.2	pegasus.condor.release	41
13.3	pegasus.condor.remove	41
13.4	pegasus.condor.output.stream	41
13.5	pegasus.condor.error.stream	42
13.6	pegasus.dagman.retry	42
13.7	pegasus.dagman.maxpre	42
13.8	pegasus.dagman.maxpost	43
13.9	pegasus.dagman.maxjobs	43
13.10	pegasus.dagman.nofity	43
13.11	pegasus.dagman.verbose	44
14	Workflow Partitioning And Job Clustering Properties	44
14.1	pegasus.partitioner.label.key	44
14.2	pegasus.partitioner.horziontal.collapse.[txname]	45
14.3	pegasus.partitioner.horziontal.bundle.[txname]	45

Contents	4
----------	---

14.4	pegasus.clusterer.job.aggregator	46
14.5	pegasus.clusterer.job.aggregator.seqexec.hasgloballog	46
14.6	pegasus.clusterer.label.key	47

15 Miscellaneous Properties	47
------------------------------------	-----------

15.1	pegasus.job.priority	47
15.2	pegasus.catalog.transformation.mapper	48
15.3	pegasus.selector.transformation	49
15.4	pegasus.auth.gridftp.timeout	49

Index	50
--------------	-----------

1 Introduction

This file is the reference guide to all properties regarding the Pegasus Workflow Planner, and their respective default values. Please refer to the user guide for a discussion when and which properties to use to configure various components. Please note that the values rely on proper capitalization, unless explicitly noted otherwise.

Some properties rely with their default on the value of other properties. As a notation, the curly braces refer to the value of the named property. For instance, `${pegasus.home}` means that the value depends on the value of the `pegasus.home` property plus any noted additions. You can use this notation to refer to other properties, though the extent of the substitutions are limited. Usually, you want to refer to a set of the standard system properties. Nesting is not allowed. Substitutions will only be done once.

There is a priority to the order of reading and evaluating properties. Usually one does not need to worry about the priorities. However, it is good to know the details of when which property applies, and how one property is able to overwrite another.

1. Property definitions in the system property file, usually found in `${pegasus.home.sysconfdir}/properties`, have the lowest priority. These properties are expected to be set up by the submit host's administrator.
2. The properties defined in the user property file `${user.home}/.pegasusrc` have higher priority. These can overwrite settings found in the system's properties. A set of sensible property values to set on a production system is shown below.
3. Commandline properties have the highest priority. Each commandline property is introduced by a `-D` argument. Note that these arguments are parsed by the shell wrapper, and thus the `-D` arguments must be the first arguments to any command. Commandline properties are useful for debugging purposes.

The following example provides a sensible set of properties to be set by the user property file. These properties use mostly non-default settings. It is an example only, and will not work for you:

```
pegasus.catalog.provenance InvocationSchema
pegasus.catalog.*.db.driver Postgres
pegasus.catalog.*.db.url jdbc:postgresql:${user.name}
pegasus.catalog.*.db.user ${user.name}
pegasus.catalog.*.db.password XXXXXXXXX
```

```

pegasus.catalog.replica RLS
pegasus.catalog.replica.url rls://smarty.isi.edu
pegasus.catalog.transformation File
pegasus.catalog.transformation.file    ${pegasus.home}/var/sample.tc.data
pegasus.catalog.site XML
pegasus.catalog.site.file                ${pegasus.home}/etc/sample.sites.xml

```

If you are in doubt which properties are actually visible, a sample application called `show-properties` dumps all properties after reading and prioritizing them.

Property : `pegasus.home` Systems : all Type : directory location string Default : `""$PEGASUS_HOME""`

The property `pegasus.home` cannot be set in the property file. Any of the shell wrapper scripts for the applications will set this property from the value of the environment variable `$PEGASUS_HOME`. Knowledge about this property is important for developers who want to invoke PEGASUS classes without the shell wrappers.

`pegasus.home ""$PEGASUS_HOME""`

2 Property Files And Locations

This section describes the property file locations. Please refer to the introduction on issues about precedence of property definitions.

2.1 pegasus.properties

Systems	all
Type	file location string
Default	<code>\${pegasus.home.sysconfdir}/properties</code>

The system-wide properties file will be looked for in its default place. It will usually reside in `$PEGASUS_HOME/etc` as file named `properties`.

2.2 pegasus.user.properties

Systems	all
Type	file location string
Default	<code>\${user.home}/.pegasusrc</code>

Each user can overwrite the system-wide properties with his or her own definitions. The user properties rely on the system's notion of the user home directory, as reflected in the JRE system properties. In the user's home directory, a file `.pegasusrc` will be taken to contain user property definitions. Note that `${user.home}` is a system property provided by the Java run-time environment (JRE).

Older version of PEGASUS used to support a dot-chimerarc file. For a while, both files are supported. However, in the presence of both files, precedence will be granted to the dot-pegasusrc file.

3 Local Directories

This section describes the GNU directory structure conventions. GNU distinguishes between architecture independent and thus sharable directories, and directories with data specific to a platform, and thus often local. It also distinguishes between frequently modified data and rarely changing data. These two axis form a space of four distinct directories.

3.1 pegasus.home.datadir

Systems	all
Type	directory location string
Default	<code>\${pegasus.home}/share</code>

The datadir directory contains broadly visible and possibly exported configuration files that rarely change. This directory is currently unused.

3.2 pegasus.home.sysconfdir

Systems	all
Type	directory location string
Default	<code>\${pegasus.home}/etc</code>

The system configuration directory contains configuration files that are specific to the machine or installation, and that rarely change. This is the directory where the XML schema definition copies are stored, and where the base pool configuration file is stored.

3.3 pegasus.home.sharedstatedir

Systems	all
Type	directory location string
Default	<code>\${pegasus.home}/com</code>

Frequently changing files that are broadly visible are stored in the shared state directory. This is currently unused.

3.4 pegasus.home.localstatedir

Systems	all
Type	directory location string
Default	\${pegasus.home}/var

Frequently changing files that are specific to a machine and/or installation are stored in the local state directory. This directory is being used for the textual transformation catalog, and the file-based replica catalog.

4 Site Directories

The site directory properties modify the behavior of remotely run jobs. In rare occasions, it may also pertain to locally run compute jobs.

4.1 pegasus.dir.exec

System	Pegasus
Since	2.0
Type	remote directory location string
Default	(no default)
Old Name	vds.dir.exec

This property modifies the remote location work directory in which all your jobs will run. If the path is relative then it is appended to the work directory (associated with the site), as specified in the site catalog. If the path is absolute then it overrides the work directory specified in the site catalog.

4.2 pegasus.dir.storage

System	Pegasus
Since	2.0
Type	remote directory location string
Default	(no default)
Old Name	vds.dir.storage

This property modifies the remote storage location on various pools. If the path is relative then it is appended to the storage mount point specified in the pool.config file. If the path is absolute then it overrides the storage mount point specified in the pool config file.

4.3 pegasus.dir.create

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	HourGlass
Value[1]	Tentacles
Default	HourGlass
Old name	vds.dir.create.mode, vds.dir.create

If the `--randomdir` option is given to the Planner at runtime, the Pegasus planner adds nodes that create the random directories at the remote pool sites, before any jobs are actually run. The two modes determine the placement of these nodes and their dependencies to the rest of the graph.

HourGlass It adds a make directory node at the top level of the graph, and all these concat to a single dummy job before branching out to the root nodes of the original/ concrete dag so far. So we introduce a classic X shape at the top of the graph. Hence the name HourGlass.

Tentacles This option places the jobs creating directories at the top of the graph. However instead of constricting it to an hour glass shape, this mode links the top node to all the relevant nodes for which the create dir job is necessary. It looks as if the node spreads its tentacles all around. This puts more load on the DAGMan because of the added dependencies but removes the restriction of the plan progressing only when all the create directory jobs have progressed on the remote pools, as is the case in the HourGlass model.

5 Schema File Location Properties

This section defines the location of XML schema files that are used to parse the various XML document instances in the PEGASUS. The schema backups in the installed file-system permit PEGASUS operations without being online.

5.1 pegasus.schema.dax

Systems	Pegasus
Since	2.0
Type	XML schema file location string
Value[0]	\${pegasus.home.sysconfdir}/dax-2.0.xsd
Default	\${pegasus.home.sysconfdir}/dax-2.0.xsd

This file is a copy of the XML schema that describes abstract DAG files that are the result of the abstract planning process, and input into any concrete planning. Providing a copy of the schema enables the parser

to use the local copy instead of reaching out to the internet, and obtaining the latest version from the GriPhyN website dynamically.

5.2 pegasus.schema.pdax

Systems	Pegasus
Since	2.0
Type	XML schema file location string
Value[0]	\${pegasus.home.sysconfdir}/pdax-2.0.xsd
Default	\${pegasus.home.sysconfdir}/pdax-2.0.xsd

This file is a copy of the XML schema that describes partition dag files that are the result of the partitioning process. Providing a copy of the schema enables the parser to use the local copy instead of reaching out to the internet, and obtaining the latest version from the GriPhyN website dynamically.

5.3 pegasus.schema.sc

Systems	Pegasus
Since	2.0
Type	XML schema file location string
Value[0]	\${pegasus.home.sysconfdir}/sc-2.0.xsd
Default	\${pegasus.home.sysconfdir}/sc-2.0.xsd
Old name	vds.schema.poolconfig, vds.schema.sc

This file is a copy of the XML schema that describes the xml description of the site catalog, that is generated as a result of using genpoolconfig command. Providing a copy of the schema enables the parser to use the local copy instead of reaching out to the internet, and obtaining the latest version from the GriPhyN website dynamically.

5.4 pegasus.schema.ivr

Systems	all
Type	XML schema file location string
Value[0]	\${pegasus.home.sysconfdir}/ivr-2.0.xsd
Default	\${pegasus.home.sysconfdir}/ivr-2.0.xsd

This file is a copy of the XML schema that describes invocation record files that are the result of the a grid launch in a remote or local site. Providing a copy of the schema enables the parser to use the local copy instead of reaching out to the internet, and obtaining the latest version from the GriPhyN website dynamically.

6 Database Drivers For All Relational Catalogs

6.1 pegasus.catalog.<catalog-name>.db.driver

System	Pegasus
Type	Java class name
Value[0]	Postgres
Value[1]	MySQL
Value[2]	SQLServer2000 (not yet implemented!)
Value[3]	Oracle (not yet implemented!)
Default	(no default)
See also	pegasus.catalog.provenance.schema, section ?? (page ??)

The database driver class is dynamically loaded, as required by the schema. Currently, only PostGreSQL 7.3 and MySQL 4.0 are supported. Their respective JDBC3 driver is provided as part and parcel of the PEGASUS.

A user may provide their own implementation, derived from `org.griphyn.vdl.dbdriver.DatabaseDriver`, to talk to a database of their choice.

For each schema in ptc and tc, a driver is instantiated separately, which has the same prefix as the schema. This may result in multiple connections to the database backend. As fallback, the schema "*" driver is attempted.

6.2 pegasus.catalog.*.db.url

System	PTC, TC, ...
Type	JDBC database URI string
Default	(no default)
Example	<code>jdbc:postgresql:\${user.name}</code>

Each database has its own string to contact the database on a given host, port, and database. Although most driver URLs allow to pass arbitrary arguments, please use the `pegasus.catalog.catalog-name.db.*` keys or `pegasus.catalog.*.db.*` to preload these arguments. THE URL IS A MANDATORY PROPERTY FOR ANY DBMS BACKEND.

```

Postgres : jdbc:postgresql:[//hostname[:port]]database
MySQL    : jdbc:mysql://[hostname[:port]]/database
SQLServer: jdbc:microsoft:sqlserver://hostname:port
Oracle   : jdbc:oracle:thin:[user/password]@//host[:port]/service

```

6.3 pegasus.catalog.*.db.user

System	PTC, TC, ...
Type	string
Default	(no default)
Example	\${user.name}

In order to access a database, you must provide the name of your account on the DBMS. This property is database-independent. THIS IS A MANDATORY PROPERTY FOR MANY DBMS BACKENDS.

6.4 pegasus.catalog.*.db.password

System	PTC, TC, ...
Type	string
Default	(no default)
Example	\${user.name}

In order to access a database, you must provide an optional password of your account on the DBMS. This property is database-independent. THIS IS A MANDATORY PROPERTY, IF YOUR DBMS BACKEND ACCOUNT REQUIRES A PASSWORD.

6.5 pegasus.catalog.*.db.*

System	PTC, TC, WORK, RC
--------	-------------------

Each database has a multitude of options to control in fine detail the further behaviour. You may want to check the JDBC3 documentation of the JDBC driver for your database for details. The keys will be passed as part of the connect properties by stripping the "pegasus.catalog.catalog-name.db." prefix from them. The catalog-name can be replaced by the following values provenance for Provenance Catalog (PTC) transformation for Transformation Catalog (TC) replica for Replica Catalog (RC) work for Workflow Catalog

Postgres 7.3 parses the following properties:

```

pegasus.catalog.*.db.user
pegasus.catalog.*.db.password
pegasus.catalog.*.db.PGHOST
pegasus.catalog.*.db.PGPORT
pegasus.catalog.*.db.charSet
pegasus.catalog.*.db.compatible

```

MySQL 4.0 parses the following properties:

```
pegasus.catalog.*.db.user
pegasus.catalog.*.db.password
pegasus.catalog.*.db.databaseName
pegasus.catalog.*.db.serverName
pegasus.catalog.*.db.portNumber
pegasus.catalog.*.db.socketFactory
pegasus.catalog.*.db.strictUpdates
pegasus.catalog.*.db.ignoreNonTxTables
pegasus.catalog.*.db.secondsBeforeRetryMaster
pegasus.catalog.*.db.queriesBeforeRetryMaster
pegasus.catalog.*.db.allowLoadLocalInfile
pegasus.catalog.*.db.continueBatchOnError
pegasus.catalog.*.db.pedantic
pegasus.catalog.*.db.useStreamLengthsInPrepStmts
pegasus.catalog.*.db.useTimezone
pegasus.catalog.*.db.relaxAutoCommit
pegasus.catalog.*.db.paranoid
pegasus.catalog.*.db.autoReconnect
pegasus.catalog.*.db.capitalizeTypeNames
pegasus.catalog.*.db.ultraDevHack
pegasus.catalog.*.db.strictFloatingPoint
pegasus.catalog.*.db.useSSL
pegasus.catalog.*.db.useCompression
pegasus.catalog.*.db.socketTimeout
pegasus.catalog.*.db.maxReconnects
pegasus.catalog.*.db.initialTimeout
pegasus.catalog.*.db.maxRows
```

```

pegasus.catalog.*.db.useHostsInPrivileges
pegasus.catalog.*.db.interactiveClient
pegasus.catalog.*.db.useUnicode
pegasus.catalog.*.db.characterEncoding

```

MS SQL Server 2000 support the following properties (keys are case-insensitive, e.g. both "user" and "User" are valid):

```

pegasus.catalog.*.db.User
pegasus.catalog.*.db.Password
pegasus.catalog.*.db.DatabaseName
pegasus.catalog.*.db.ServerName
pegasus.catalog.*.db.HostProcess
pegasus.catalog.*.db.NetAddress
pegasus.catalog.*.db.PortNumber
pegasus.catalog.*.db.ProgramName
pegasus.catalog.*.db.SendStringParametersAsUnicode
pegasus.catalog.*.db.SelectMethod

```

6.6 pegasus.catalog.provenance

System	Provenance Tracking Catalog (PTC)
Type	Java class name
Value[0]	InvocationSchema
Value[1]	NXDInvSchema
Default	(no default)
See also	pegasus.catalog.*.db.driver, section ?? (page ??)

This property denotes the schema that is being used to access a PTC. The PTC is usually not a standard installation. If you use a database backend, you most likely have a schema that supports PTCs. By default, no PTC will be used.

Currently only the InvocationSchema is available for storing the provenance tracking records. Beware, this can become a lot of data. The values are names of Java classes. If no absolute Java classname is given, "org.griphyn.vdl.dbschema." is prepended. Thus, by deriving from the DatabaseSchema API, and implementing the PTC interface, users can provide their own classes here.

Alternatively, if you use a native XML database like eXist, you can store data using the NXDInvSchema. This will avoid using any of the other database driver properties.

6.7 pegasus.catalog.work

System	Work Catalog
Type	Java class name
Value[0]	Database
Default	Database
See also	pegasus.catalog.*.db.driver, section ?? (page ??)

This property denotes the schema that is being used to store workflow monitoring entries. If the catalog

7 Catalog Properties

7.1 Replica Catalog

7.1.1 pegasus.catalog.replica

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	RLS
Value[1]	LRC
Value[2]	JDBCRC
Value[3]	SimpleFile
Default	RLS
Old name	vds.rc, vds.replica.mode

Pegasus queries a Replica Catalog to discover the physical filenames (PFN) for input files specified in the DAX. Pegasus can interface with various types of Replica Catalogs. This property specifies which type of Replica Catalog to use during the planning process.

RLS RLS (Replica Location Service) is a distributed replica catalog, which ships with GT4. There is an index service called Replica Location Index (RLI) to which 1 or more Local Replica Catalog (LRC) report. Each LRC can contain all or a subset of mappings. In this mode, Pegasus queries the central RLI to discover in which LRC's the mappings for a LFN reside. It then queries the individual LRC's for the PFN's. To use RLS, the user additionally needs to set the property `pegasus.catalog.replica.url` to specify the URL for the RLI to query. Details about RLS can be found at <http://www.globus.org/toolkit/data/rls/>

LRC If the user does not want to query the RLI, but directly a single Local Replica Catalog. To use LRC, the user additionally needs to set the property `pegasus.catalog.replica.url` to specify the URL for the LRC to query. Details about RLS can be found at <http://www.globus.org/toolkit/data/rls/>

JDBCRC In this mode, Pegasus queries a SQL based replica catalog that is accessed via JDBC. The sql schema's for this catalog can be found at `$PEGASUS_HOME/sql` directory. To use JDBCRC, the user additionally needs to set the following properties

1. `pegasus.catalog.replica.db.url`
2. `pegasus.catalog.replica.db.user`
3. `pegasus.catalog.replica.db.password`

SimpleFile In this mode, Pegasus queries a file based replica catalog. It is neither transactionally safe, nor advised to use for production purposes in any way. Multiple concurrent instances **will clobber** each other!p. The site attribute should be specified whenever possible. The attribute key for the site attribute is "pool".

The LFN may or may not be quoted. If it contains linear whitespace, quotes, backslash or an equality sign, it must be quoted and escaped. Ditto for the PFN. The attribute key-value pairs are separated by an equality sign without any whitespaces. The value may be in quoted. The LFN sentiments about quoting apply.

```
LFN PFN
```

```
LFN PFN a=b [..]
```

```
LFN PFN a="b" [..]
```

```
"LFN w/LWS" "PFN w/LWS" [..]
```

To use SimpleFile, the user additionally needs to specify `pegasus.catalog.replica.file` property to specify the path to the file based RC.

7.1.2 `pegasus.catalog.replica.url`

System	Pegasus
Since	2.0
Type	URI string
Default	(no default)
Old name	<code>vds.rc.url</code> , <code>vds.rls.url</code>

When using the modern RLS replica catalog, the URI to the Replica catalog must be provided to Pegasus to enable it to look up filenames. There is no default.

7.1.3 pegasus.catalog.replica.chunk.size

System	Pegasus, rc-client
Since	2.0
Type	Integer
Default	1000
Old name	vds.rc.chunk.size

The rc-client takes in an input file containing the mappings upon which to work. This property determines, the number of lines that are read in at a time, and worked upon at together. This allows the various operations like insert, delete happen in bulk if the underlying replica implementation supports it.

7.1.4 pegasus.catalog.replica.lrc.ignore

System	Replica Catalog - RLS
Since	2.0
Type	comma separated list of LRC urls
Default	(no default)
See also	pegasus.catalog.replica.lrc.restrict, section 7.1.5 (page 17)
Old name	vds.rc.lrc.ignore

Certain users may like to skip some LRCs while querying for the physical locations of a file. If some LRCs need to be skipped from those found in the rli then use this property. You can define either the full URL or partial domain names that need to be skipped. E.g. If a user wants rls://smarty.isi.edu and all LRCs on usc.edu to be skipped then the property will be set as pegasus.rls.lrc.ignore=rls://smarty.isi.edu,usc.edu

7.1.5 pegasus.catalog.replica.lrc.restrict

System	Replica Catalog - RLS
Since	1.3.9
Type	comma separated list of LRC urls
Default	(no default)
See also	pegasus.catalog.replica.lrc.ignore, section 7.1.4 (page 17)
Old name	vds.rc.lrc.ignore

This property applies a tighter restriction on the results returned from the LRCs specified. Only those PFNs are returned that have a pool attribute associated with them. The property "pegasus.rc.lrc.ignore" has a higher priority than "pegasus.rc.lrc.restrict". For example, in case a LRC is specified in both properties, the LRC would be ignored (i.e. not queried at all instead of applying a tighter restriction on the results returned).

7.1.6 pegasus.catalog.replica.cache.asrc

System	Pegasus
Since	2.0
Type	Boolean
Value[0]	false
Value[1]	true
Default	false
See also	pegasus.catalog.replica, section 7.1.1 (page 15)
Old name	vds.cache.asrc

This property determines whether to treat the cache file specified as a supplemental replica catalog or not. User can specify on the command line to gencdag a comma separated list of cache files using the `–cache` option. By default, the LFN->PFN mappings contained in the cache file are treated as cache, i.e if an entry is found in a cache file the replica catalog is not queried. This results in only the entry specified in the cache file to be available for replica selection.

Setting this property to true, results in the cache files to be treated as supplemental replica catalogs. This results in the mappings found in the replica catalog (as specified by `pegasus.catalog.replica`) to be merged with the ones found in the cache files. Thus, mappings for a particular LFN found in both the cache and the replica catalog are available for replica selection.

7.2 Site Catalog

7.2.1 pegasus.catalog.site

System	Site Catalog
Since	2.0
Type	enumeration
Value[0]	XML
Value[1]	Text
Default	XML
Old name	vds.pool.mode, vds.sc

The site catalog file is available in two major flavors:

1. The "XML" format is an XML-based file. It is generated using the `genpoolconfig` client application program that is shipped with Pegasus. The XML input file for Pegasus can be generated in various ways, that can be used exclusively or combined at your option:
 - a) It can also be published by converting the new, easier to read and modify local multiline pool config file. An example is provided in `sample.pool.config.new`. Use this option if you have no network connectivity, or for tests.

2. The "Text" format is a multiline site catalog format. It is described in the site catalog guide. It can be directly given to Pegasus starting with PEGASUS-1.4

7.2.2 pegasus.catalog.site.file

System	Site Catalog
Since	2.0
Type	file location string
Default	\${pegasus.home.sysconfdir}/sites.xml \${pegasus.home.sysconfdir}/sites.txt
See also	pegasus.catalog.site, section 7.2.1 (page 18)
Old name	vds.pool.file, vds.sc

Running things on the grid requires an extensive description of the capabilities of each compute cluster, commonly termed "site". This property describes the location of the file that contains such a site description. As the format is currently in flow, please refer to the userguide and Pegasus for details which format is expected. The default value is dependant on the value specified for the property pegasus.sc . pegasus.sc denotes the type of site catalog being used.

7.3 Transformation Catalog

7.3.1 pegasus.catalog.transformation

System	Transformation Catalog
Since	2.0
Type	enumeration
Value[0]	File
Value[1]	Database
Default	File
See also	pegasus.catalog.*.driver, section ?? (page ??)
See also	pegasus.catalog.transformation.file, section 7.3.2 (page 20)
Old name	vds.tc.mode, vds.tc

File In this mode, a file format is understood. The file is read and cached in memory. Any modifications, as adding or deleting, causes an update of the memory and hence to the file underneath. All queries are done against the memory representation. The new TC file format uses 6 columns:

1. The resource ID is represented in the first column.
2. The logical transformation uses the colonized format ns::name:vs.
3. The path to the application on the system

4. The installation type is identified by one of the following keywords - all upper case: `INSTALLED`, `STATIC_BINARY`, `DYNAMIC_BINARY`, `SCRIPT`. If not specified, or `NULL` is used, the type defaults to `INSTALLED`.
5. The system is of the format `ARCH::OS[:VER:GLIBC]`. The following arch types are understood: `"INTEL32"`, `"INTEL64"`, `"SPARCV7"`, `"SPARCV9"`. The following os types are understood: `"LINUX"`, `"SUNOS"`, `"AIX"`. If unset or `NULL`, defaults to `INTEL32::LINUX`.
6. Profiles are written in the format `NS::KEY=VALUE,KEY2=VALUE;NS2::KEY3=VALUE3`. Multiple key-values for same namespace are separated by a comma `,` and multiple namespaces are separated by a semicolon `;`. If any of your profile values contains a comma you must not use the namespace abbreviator.

Database In this mode, the transformation catalog is kept in a relational database. Currently mysql DB and Postgre are supported. To set up the the database, use the schema in `$PEGASUS_HOME/sql/create-my-init.sql` followed by `$PEGASUS_HOME/sql/create-my-tc.sql` .

```
The following properties need to be set

pegasus.catalog.transformation.db = MySQL|Postgres

pegasus.catalog.transformation.db.url =
jdbc:mysql://[hostname[:port]]/database |
jdbc:postgres://[hostname[:port]]/database

pegasus.catalog.transformation.db.username = dbusername

pegasus.catalog.transformation.db.password = password
```

If the `pegasus.catalog.transformation.db.*` properties are not defined, the database implementation picks up the properties specified by `pegasus.catalog.*.db.*` .

Future modifications to the TC may extend the enumeration. To implement your own TC implementation see `org.girphyn.cPlanner.tc.TCMechanism`. To load the class set `pegasus.catalog.transformation` to the TC implementation class.

7.3.2 pegasus.catalog.transformation.file

Systems	Transformation Catalog
Type	file location string
Default	<code>\${pegasus.home.localstatedir}/tc.data</code>
See also	<code>pegasus.catalog.transformation</code> , section 7.3.1 (page 19)
Old name	<code>vds.tc.file</code>

The transformation catalog is a 6 column textual file that describes in a simple column based format the mapping from a logical transformation for each pool to the physical application, and optional environment settings. All concrete planners (Pegasus and Euryale) use this repository to map the ITR from the abstract DAX into an application invocation. Refer to the user guide for details.

8 Replica Selection Properties

8.1 pegasus.selector.replica

System	Replica Selection
Since	2.0
Type	URI string
Default	default
See also	pegasus.replica.*.ignore.stagein.sites, section ?? (page ??)
See also	pegasus.replica.*.prefer.stagein.sites, section ?? (page ??)
Old name	vds.rc.selector, vds.replica.selector

Each job in the DAX maybe associated with input LFN's denoting the files that are required for the job to run. To determine the physical replica (PFN) for a LFN, Pegasus queries the replica catalog to get all the PFN's (replicas) associated with a LFN. Pegasus then calls out to a replica selector to select a replica amongst the various replicas returned. This property determines the replica selector to use for selecting the replicas.

Default If a PFN that is a file URL (starting with file:///) and has a pool attribute matching to the site handle of the site where the compute is to be run is found, then that is returned. Else, a random PFN is selected amongst all the PFN's that have a pool attribute matching to the site handle of the site where a compute job is to be run. Else, a random pfn is selected amongst all the PFN's.

Restricted This replica selector, allows the user to specify good sites and bad sites for staging in data to a particular compute site. A good site for a compute site X, is a preferred site from which replicas should be staged to site X. If there are more than one good sites having a particular replica, then a random site is selected amongst these preferred sites.

A bad site for a compute site X, is a site from which replica's should not be staged. The reason of not accessing replica from a bad site can vary from the link being down, to the user not having permissions on that site's data.

The good | bad sites are specified by the properties

```
pegasus.replica.*.prefer.stagein.sites
pegasus.replica.*.ignore.stagein.sites
```

where the * in the property name denotes the name of the compute site. A * in the property key is taken to mean all sites.

The `pegasus.replica.*.prefer.stagein.sites` property takes precedence over `pegasus.replica.*.ignore.stagein.sites` property i.e. if for a site X, a site Y is specified both in the ignored and the preferred set, then site Y is taken to mean as only a preferred site for a site X.

8.2 `pegasus.selector.replica.*.ignore.stagein.sites`

System	Replica Selection
Type	comma separated list of sites
Since	2.0
Default	no default
See also	<code>pegasus.selector.replica</code> , section 8.1 (page 21)
See also	<code>pegasus.selector.replica.*.prefer.stagein.sites</code> , section 8.3 (page 22)
Old name	<code>vds.rc.restricted.sites</code> , <code>vds.replica.*.ignore.stagein.sites</code>

A comma separated list of storage sites from which to never stage in data to a compute site. The property can apply to all or a single compute site, depending on how the * in the property name is expanded.

The * in the property name means all compute sites unless replaced by a site name.

For e.g setting `pegasus.selector.replica.*.ignore.stagein.sites` to `usc` means that ignore all replicas from site `usc` for staging in to any compute site. Setting `pegasus.replica.isi.ignore.stagein.sites` to `usc` means that ignore all replicas from site `usc` for staging in data to site `isi`.

8.3 `pegasus.selector.replica.*.prefer.stagein.sites`

System	Replica Selection
Type	comma separated list of sites
Since	2.0
Default	no default
See also	<code>pegasus.selector.replica</code> , section 8.1 (page 21)
See also	<code>pegasus.selector.replica.*.ignore.stagein.sites</code> , section 8.2 (page 22)
Old name	<code>vds.replica.*.prefer.stagein.sites</code>

A comma separated list of preferred storage sites from which to stage in data to a compute site. The property can apply to all or a single compute site, depending on how the * in the property name is expanded.

The * in the property name means all compute sites unless replaced by a site name.

For e.g setting `pegasus.selector.replica.*.prefer.stagein.sites` to `usc` means that prefer all replicas from site `usc` for staging in to any compute site. Setting `pegasus.replica.isi.prefer.stagein.sites` to `usc` means that prefer all replicas from site `usc` for staging in data to site `isi`.

9 Site Selection Properties

9.1 pegasus.selector.site

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	Random
Value[1]	RoundRobin
Value[2]	NonJavaCallout
Value[3]	Group
Default	Random
See also	pegasus.selector.site.path, section 9.2 (page 24)
See also	pegasus.selector.site.timeout, section 9.4 (page 25)
See also	pegasus.selector.site.keep.tmp, section 9.5 (page 25)
See also	pegasus.selector.site.env.*, section ?? (page ??)
Old name	vds.site.selector

The site selection in Pegasus can be on basis of any of the following strategies.

Random In this mode, the jobs will be randomly distributed among the sites that can execute them.

RoundRobin In this mode, the jobs will be assigned in a round robin manner amongst the sites that can execute them. Since each site cannot execute everytype of job, the round robin scheduling is done per level on a sorted list. The sorting is on the basis of the number of jobs a particular site has been assigned in that level so far. If a job cannot be run on the first site in the queue (due to no matching entry in the transformation catalog for the transformation referred to by the job), it goes to the next one and so on. This implementation defaults to classic round robin in the case where all the jobs in the workflow can run on all the sites.

NonJavaCallout In this mode, Pegasus will callout to an external site selector. In this mode a temporary file is prepared containing the job information that is passed to the site selector as an argument while invoking it. The path to the site selector is specified by setting the property `pegasus.site.selector.path`. The environment variables that need to be set to run the site selector can be specified using the properties with a `pegasus.site.selector.env.` prefix. The temporary file contains information about the job that needs to be scheduled. It contains key value pairs with each key value pair being on a new line and separated by a `=`.

The following pairs are currently generated for the site selector temporary file that is generated in the NonJavaCallout.

version	is the version of the site selector api,currently 2.0.
transformation	is the fully-qualified definition identifier for the transformation (TR) namespace::name:version.
derivation	is teh fully qualified definition identifier for the derivation (DV), namespace::name:version.
job.level	is the job's depth in the tree of the workflow DAG.
job.id	is the job's ID, as used in the DAX file.
resource.id	is a pool handle, followed by whitespace, followed by a gridftp server. Typically, each gridftp server is enumerated once, so you may have multiple occurances of the same site. There can be multiple occurances of this key.
input.lfn	is an input LFN, optionally followed by a whitespace and file size. There can be multiple occurances of this key,one for each input LFN required by the job.
wf.name	label of the dax, as found in the DAX's root element. wf.index is the DAX index, that is incremented for each partition in case of deferred planning.
wf.time	is the mtime of the workflow.
wf.manager	is the name of the workflow manager being used .e.g condor
vo.name	is the name of the virtual organization that is running this workflow. It is currently set to NONE
vo.group	unused at present and is set to NONE.

Group In this mode, a group of jobs will be assigned to the same site that can execute them. The use of the PEGASUS profile key group in the dax, associates a job with a particular group. The jobs that do not have the profile key associated with them, will be put in the defaultMANAGEMENT group. The jobs in the default group are handed over to the "Random" Site Selector for scheduling.

9.2 pegasus.selector.site.path

System	Site Selector
Since	2.0
Type	String
Old name	vds.site.selector.path

If one calls out to an external site selector using the NonJavaCallout mode, this refers to the path where the site selector is installed. In case other strategies are used it does not need to be set.

9.3 pegasus.site.selector.env.*

System	Pegasus
Since	1.2.3
Type	String

The environment variables that need to be set while callout to the site selector. These are the variables that the user would set if running the site selector on the command line. The name of the environment variable is got by stripping the keys of the prefix "pegasus.site.selector.env." prefix from them. The value of the environment variable is the value of the property.

e.g pegasus.site.selector.path.LD_LIBRARY_PATH /globus/lib would lead to the site selector being called with the LD_LIBRARY_PATH set to /globus/lib.

9.4 pegasus.selector.site.timeout

System	Site Selector
Since	2.0
Type	non negative integer
Default	60
Old name	vds.site.selector.timeout

It sets the number of seconds Pegasus waits to hear back from an external site selector using the NonJava-Callout interface before timing out.

9.5 pegasus.selector.site.keep.tmp

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	onerror
Value[1]	always
Value[2]	never
Default	onerror
Old name	vds.site.selector.keep.tmp

It determines whether Pegasus deletes the temporary input files that are generated in the temp directory or not. These temporary input files are passed as input to the external site selectors.

A temporary input file is created for each that needs to be scheduled.

10 Transfer Configuration Properties

10.1 pegasus.transfer.*.impl

System	Pegasus
Type	enumeration
Value[0]	OldGUC
Value[1]	Transfer
Value[2]	T2
Value[3]	Stork
Value[4]	RFT
Value[5]	CRFT
Value[6]	SRM
Default	Transfer
Old name	vds.transfer.mode, vds.transfer, vds.transfer.*.impl
See also	pegasus.transfer.refiner, section 10.2 (page 28)
See also	pegasus.transfer.single.quote, section 10.6 (page 30)
Since	2.0

Each compute job usually has data products that are required to be staged in to the execution site, materialized data products staged out to a final resting place, or staged to another job running at a different site. This property determines the underlying grid transfer tool that is used to manage the transfers.

The * in the property name can be replaced to achieve finer grained control to dictate what type of transfer jobs need to be managed with which grid transfer tool.

Usually, the arguments with which the client is invoked can be specified by

- the property `pegasus.transfer.arguments`
- associating the PEGASUS profile key `transfer.arguments`

The table below illustrates all the possible variations of the property.

Property Name	Applies to
<code>pegasus.transfer.stagein.impl</code>	the stage in transfer jobs
<code>pegasus.transfer.stageout.impl</code>	the stage out transfer jobs
<code>pegasus.transfer.inter.impl</code>	the inter pool transfer jobs
<code>pegasus.transfer.*.impl</code>	apply to types of transfer jobs

The various grid transfer tools that can be used to manage data transfers are explained below

OldGUC This refers to the old guc client that can manage only one file transfer per invocation. This should only be used if you have a very old globus installation.

There should be an entry in the transformation catalog with the fully qualified name as globus-url-copy for all the sites where workflow is run, or on the local site in case of third party transfers.

Transfer This is a wrapper around the globus-url-copy client that is distributed with the PEGASUS. It is able to manage multiple file transfers at the same time. For more details refer to the manpage of transfer. Transfer binary can be found at \$PEGASUS_HOME/bin/transfer.

There should be an entry in the transformation catalog with the fully qualified name as transfer for all the sites where workflow is run, or on the local site in case of third party transfers.

T2 This is a successor to the Transfer tool. It is also distributed with the PEGASUS and the binary can be found at \$PEGASUS_HOME/bin/T2. In addition, T2 can handle conditional transfers, whereby it signals a success in case of failures during transfers. The input file for the transfer job that is constructed contains multiple source and destination urls for the same transfer. The transfer fails only if all pair candidates are attempted without success.

There should be an entry in the transformation catalog with the fully qualified name as pegasus::T2 for all the sites where workflow is run, or on the local site in case of third party transfers.

Stork This results in the transfers being scheduled using Stork, a data placement scheduler out of Condor. It leads to the transfer submit files generated in stork format. At present it can handle only one file transfer per invocation like OldGUC.

RFT RFT is a GT4 based webservice, that does reliable file transfer. This refers to using the java based rft-client distributed with GLOBUS to do the file transfers.

The RFT transfer implementation parses the following properties.

Name	Default Value
pegasus.transfer.rft.host	localhost
pegasus.transfer.rft.port	8080
pegasus.transfer.rft.binary	true
pegasus.transfer.rft.bs	16000
pegasus.transfer.rft.tcpbs	16000
pegasus.transfer.rft.notpt	false
pegasus.transfer.rft.streams	1
pegasus.transfer.rft.DCAU	true
pegasus.transfer.rft.processes	1
pegasus.transfer.rft.retry	3

There should be an entry in the transformation catalog with the fully qualified name as globus::rft on the local site.

CRFT RFT is a GT4 based webservice, that does reliable file transfer. This refers to using the C based rft-client distributed with GLOBUS to do the file transfers.

There should be an entry in the transformation catalog with the fully qualified name as globus::crft for all the sites where workflow is run, or on the local site in case of third party transfers.

SRM This refers to srmcp client that can talk to a SRM server. The srmcp client against which it was tested is version 1.2.0 out of FNAL. At present it can handle only one file transfer per invocation. Hence, `pegasus.transfer.refiner` needs to be set to either `SDefault` or `SChain`.

There should be an entry in the transformation catalog with the fully qualified name as `SRM::srmcp` for all the sites where workflow is run, or on the local site in case of third party transfers.

10.2 pegasus.transfer.refiner

System	Pegasus
Type	enumeration
Value[0]	SDefault
Value[1]	Default
Value[2]	Bundle
Value[3]	Chain
Default	Default
Since	2.0
Old name	<code>vds.transfer</code> , <code>vds.transfer.*.impl</code>
See also	<code>pegasus.transfer.*.impl</code> , section 10.1 (page 26)
See also	<code>pegasus.transfer.single.quote</code> , section 10.6 (page 30)

This property determines how the transfer nodes are added to the workflow. The various refiners differ in the how they link the various transfer jobs, and the number of transfer jobs that are created per compute jobs.

SDefault This is the default refinement strategy for the transfer tools that can handle only one file transfer per invocation like OldGUC and Stork. Thus, each file that needs to be transferred will generate one transfer job. This takes care of file clobbering while staging in data to a remote grid site. File Clobbering can occur when two jobs scheduled on the same site require the same input file to be staged.

Default This is the default refinement strategy for the transfer tools that can handle multiple file transfers per invocation. In this, all files required by a particular transfer job will be attempted to be transferred with just one transfer job. This also takes care of file clobbering while staging in data to a remote grid site. File Clobbering can occur when two jobs scheduled on the same site require the same input file to be staged.

Bundle In this refinement strategy, the number of stage in transfer nodes that are constructed per execution site can vary. The number of transfer nodes can be specified, by associating the pegasus profile "bundle.stagein". The profile can either be associated with the execution site in the site catalog or with the "transfer" executable in the transformation catalog. The value in the transformation catalog overrides the one in the site catalog. This refinement strategy extends from the Default refiner, and thus takes care of file clobbering while staging in data.

Chain In this refinement strategy, chains of stagein transfer nodes are constructed. A chain means that the jobs are sequentially dependant upon each other i.e. at any moment, only one stage in transfer job will run per chain. The number of chains can be specified by associating the pegasus profile "chain.stagein". The profile can either be associated with the execution site in the site catalog or with the "transfer" executable in the transformation catalog. The value in the transformation catalog overrides the one in the site catalog. This refinement strategy extends from the Default refiner, and thus takes care of file clobbering while staging in data.

10.3 pegasus.transfer.arguments

System	Pegasus
Since	2.0
Type	String
Default	(no default)
Old name	vds.transfer.arguments

This determines the arguments with which the transfer executable is invoked. The transfer executable that is invoked is dependant upon the transfer mode that has been selected. The property can be overloaded by associated the pegasus profile key transfer.arguments either with the site in the site catalog or the corresponding transfer executable in the transformation catalog.

10.4 pegasus.transfer.links

System	Pegasus
Type	boolean
Default	false
Since	2.0
Old name	vds.transfer.mode.links, vds.transfer.links
See also	pegasus.transfer, section ?? (page ??)
See also	pegasus.transfer.force, section 10.5 (page 30)

If this is set, and the transfer implementation is set to Transfer i.e. using the transfer executable distributed with the PEGASUS. On setting this property, if Pegasus while fetching data from the RLS sees a pool attribute associated with the PFN that matches the execution pool on which the data has to be transferred to, Pegasus instead of the URL returned by the RLS replaces it with a file based URL. This supposes that the if the pools match the filesystems are visible to the remote execution directory where input data resides. On seeing both the source and destination urls as file based URLs the transfer executable spawns a job that creates a symbolic link by calling `ln -s` on the remote pool. This ends up bypassing the GridFTP server and reduces the load on it, and is much faster.

10.5 pegasus.transfer.force

System	Pegasus
Since	2.0
Type	boolean
Default	false
Old name	vds.transfer.force, vds.transfer.mode.force

This determines if the force option is to be passed to the underlying transfer mechanism that is being invoked. The force option is set only if the underlying grid transfer executable exposes the force option. The force option generally is -f and is supported by globus-url-copy (Single Transfer Mode), transfer (Multiple Transfer Mode), T2 (T2 transfer mode).

10.6 pegasus.transfer.single.quote

System	Pegasus
Type	boolean
Default	true
Since	2.0
Old name	vds.transfer.single.quote

This affects only the single transfer mode and all its extensions. It determines whether to introduce quoting around the URLs or not while writing out them in the condor submit files. In certain setups this alleviates the condor quoting problem.

10.7 pegasus.transfer.throttle.processes

System	Pegasus
Since	2.0
Type	integer
Default	4
See also	pegasus.transfer, section ?? (page ??)
See also	pegasus.transfer.throttle.streams, section 10.8 (page 31)
Old name	vds.transfer.throttle.processes

This property is picked up when transfer mode (pegasus.transfer) is multiple. In this mode, multiple files are transferred using a transfer executable that comes with the PEGASUS system. This transfer executable attempts to transfer multiple files by spawning multiple g-u-c processes. By default a maximum of 4 processes are spawned to transfer the files. Using this one can change the number of processes that are spawned by the transfer executable.

10.8 pegasus.transfer.throttle.streams

System	Pegasus
Since	2.0
Type	integer
Default	1
See also	pegasus.transfer, section ?? (page ??)
See also	pegasus.transfer.throttle.processes, section 10.7 (page 30)
Old name	vds.transfer.throttle.streams

Whatever the transfer mode specified, at present each uses globus-url-copy (g-u-c) as the underlying transfer mechanism. g-u-c can open multiple streams to do the ftp data transfer. This property can be used to set the number of streams that are used to transfer one file by underlying g-u-c. It directly maps to the -p option of g-u-c.

10.9 pegasus.transfer.*.thirdparty.sites

System	Pegasus
Type	comma separated list of sites
Default	no default
Old name	vds.transfer.thirdparty.pools, vds.transfer.*thirdparty.sites
Since	2.0

By default Pegasus employs the push/pull model of transfer for transferring files in and out of a site. It does use the third party transfer option that grid ftp servers provide. This list specifies the list of sites on which the user wants third party transfers instead of the normal mode. Normally for a push transfer the source URL is file://blah and destination URL is gsiftp://blah. However in case of a third party site both the urls would be gsiftp://blah. For all these sites, the transfers are actually scheduled on the submit host (site local) in the scheduler universe. The * in the property name can be replaced to achieve finer grained control to dictate what type of transfer jobs need to be third party enabled for a particular site.

The table below illustrates all the possible variations of the property.

Property Name	Applies to
pegasus.transfer.stagein.thirdparty.sites	the stage in transfer jobs
pegasus.transfer.stageout.thirdparty.sites	the stage out transfer jobs
pegasus.transfer.inter.thirdparty.sites	the inter pool transfer jobs
pegasus.transfer.*.thirdparty.sites	apply to types of transfer jobs

In addition * can be specified as a property value, to designate that it applies to all sites.

10.10 pegasus.transfer.*.thirdparty.remote

System	Pegasus
Type	comma separated list of sites
Default	no default
Since	2.0
Old name	vds.transfer.*.thirdparty.remote

By default Pegasus schedules all the third party transfers on the submit host i.e. site 'local' in the site catalog. This property, allows the user to run third party transfers on the remote sites, instead of site 'local'. This is useful in case of transfer tools like RFT and CRFT clients that can only do third party transfers. Hence, the default push and pull model of transfer does not work for them where one of the URL's is a file:/// URL.

The table below illustrates all the possible variations of the property.

Property Name	Applies to
pegasus.transfer.stagein.thirdparty.remote	the stage in transfer jobs
pegasus.transfer.stageout.thirdparty.remote	the stage out transfer jobs
pegasus.transfer.inter.thirdparty.remote	the inter pool transfer jobs
pegasus.transfer.*.thirdparty.remote	apply to types of transfer jobs

In addition * can be specified as a property value, to designate that it applies to all sites.

10.11 pegasus.transfer.staging.delimiter

System	Pegasus
Since	2.0
Type	String
Default	:
See also	pegasus.transformation.selector, section ?? (page ??)
Old name	vds.transfer.staging.delimiter

Pegasus supports executable staging as part of the workflow. Currently staging of statically linked executables is supported only. An executable is normally staged to the work directory for the workflow/partition on the remote site. The basename of the staged executable is derived from the namespace,name and version of the transformation in the transformation catalog. This property sets the delimiter that is used for the construction of the name of the staged executable.

10.12 `pegasus.transfer.disable.chmod.sites`

System	Pegasus
Since	2.0
Type	comma separated list of sites
Default	no default
Old name	vds.transfer.disable.chmod.sites

During staging of executables to remote sites, chmod jobs are added to the workflow. These jobs run on the remote sites and do a chmod on the staged executable. For some sites, this may not be required. The permissions might be preserved, or there may be an automatic mechanism that does it.

This property allows you to specify the list of sites, where you do not want the chmod jobs to be executed. For those sites, the chmod jobs are replaced by NoOP jobs. The NoOP jobs are executed by Condor, and instead will immediately have a terminate event written to the job log file and removed from the queue.

10.13 `pegasus.transfer.proxy`

System	Pegasus
Since	2.0
Type	Boolean
Default	false
Old name	vds.transfer.proxy

By default, CondorG transfers a limited proxy to the remote site, while running jobs in the grid universe. However, certain grid ftp servers (like those in front of SRB) require a fully user proxy. In this case, the planners need to transfer the proxy along with the transfer job using Condor file transfer mechanisms. This property triggers Pegasus into creating the appropriate condor commands, that transfer the proxy from the submit host to the remote host. The source location is determined from the value of the X509_USER_KEY env profile key, that is associated with site local in the site catalog.

10.14 `pegasus.transfer.*.priority`

System	Pegasus
Type	Integer
Default	no default
Since	2.0
See also	pegasus.job.priority, section 15.1 (page 47)
Old name	vds.transfer.*.priority

This property sets the priority of the transfer jobs. It results in the creation of a priority classad in the submit files for the transfer jobs. To give priorities to different types of transfer jobs you can specify the following more specific properties.

Property Name	Applies to
pegasus.transfer.stagein.priority	the stage in transfer jobs
pegasus.transfer.stageout.priority	the stage out transfer jobs
pegasus.transfer.inter.priority	the inter pool transfer jobs

11 Gridstart And Exitcode Properties

11.1 pegasus.gridstart

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	Kickstart
Value[1]	NoGridStart
Default	Kickstart
Old Name	vds.gridstart

Jobs that are launched on the grid maybe wrapped in a wrapper executable/script that enables information about about the execution, resource consumption, and - most importantly - the exitcode of the remote application. At present, a job scheduled on a remote site is launched with a gridstart if site catalog has the corresponding gridlaunch attribute set and the job being launched is not MPI.

Kickstart In this mode, all the jobs are lauched via kickstart. The kickstart executable is a light-weight program which connects the stdin,stdout and stderr filehandles for PEGASUS jobs on the remote site. Kickstart is an executable distributed with PEGASUS that can generally be found at `${pegasus.home.bin}/kickstart`.

NoGridStart In this mode, all the jobs are launched directly on the remote site. Each job's stdin,stdout and stderr are connected to condor commands in a manner to ensure that they are sent back to the submit host.

Support for a new gridstart (K2) is expected to be added soon.

11.2 pegasus.gridstart.invoke.always

System	Pegasus
Since	2.0
Type	Boolean
Default	false
See also	pegasus.gridstart.invoke.length, section 11.3 (page 35)
Old Name	vds.gridstart.invoke.always

Condor has a limit in it, that restricts the length of arguments to an executable to 4K. To get around this limit, you can trigger Kickstart to be invoked with the -I option. In this case, an arguments file is prepared per job that is transferred to the remote end via the Condor file transfer mechanism. This way the arguments to the executable are not specified in the condor submit file for the job. This property specifies whether you want to use the invoke option always for all jobs, or want it to be triggered only when the argument string is determined to be greater than a certain limit.

11.3 pegasus.gridstart.invoke.length

System	Pegasus
Since	2.0
Type	Long
Default	4000
See also	pegasus.gridstart.invoke.always, section 11.2 (page 34)
Old Name	vds.gridstart.invoke.length

Gridstart is automatically invoked with the -I option, if it is determined that the length of the arguments to be specified is going to be greater than a certain limit. By default this limit is set to 4K. However, it can be overridden by specifying this property.

11.4 pegasus.exitcode.scope

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	all
Value[1]	none
Value[2]	essential
Default	none
Old name	vds.exitcode, vds.exitcode.mode
See also	pegasus.exitcode.arguments, section 11.7 (page 38)

Jobs that are "kickstarted" by a grid launcher report back information about the execution, resource consumption, and - most importantly - the exit code of the remote application. Armed with this knowledge, it is possible to have DAGMan stop the workflow and create a rescue workflow on remote execution errors.

1. In "all" mode, each kickstarted job's invocation record will be parsed by a DAGMan postscript.
2. In "none" mode, the default, remote failures will not abort the DAG.
3. In "essential" mode, only certain classes of remote jobs get the ability to abort the workflow, while other, non-essential, classes of jobs (at present the replica registration jobs) will not have their invocation record abort the workflow.

11.5 pegasus.exitcode.impl

System	Pegasus
Since	2.0
Type	enumeration
Value[0]	exitpost
Value[1]	exitcode
Value[2]	none
Value[3]	any legal identifier (concretely [_A-Za-z][_A-Za-z0-9]*)
Default	no default
See also	pegasus.exitcode.path.[value], section 11.6 (page 37)
See also	pegasus.exitcode.arguments, section 11.7 (page 38)
Old Name	vds.exitcode.impl

Each of the kickstarted jobs can have a postscript associated with them. This postscript is executed on the submit host, after a job has finished execution. By default, each of the GridStart implementations have their own postscripts, that are able to parse their output and determine the exit status. For example, for kickstart PEGASUS provides an exitpost utility that parses the kickstart record and can populate the data into a provenance tracking catalog. However, you can specify your own postscript to be executed if desired.

This property applies to all the jobs in the workflow. The value can be overridden selectively for jobs by associating Dagman profile key POST with the jobs either in DAX, site in the Site Catalog or the transformation in the transformation catalog.

The path to the postscript can be specified by the property pegasus.exitcode.path.[value] and arguments by pegasus.exitcode.arguments, where [value] is the value of this property (pegasus.exitcode.impl).

The following types of postscript are supported.

exitpost This is the postscript that can be found at \$PEGASUS_HOME/bin/exitpost, and is by default invoked on all jobs that are launched via kickstart. It is a wrapper around exitcode, that makes backups of the processed files. It assumes that the error file ends in ".err" (to be fixed), and looks into it for certain PBS failures like wall-time exceeded, which may fail a job without producing good kickstart output. This thin wrapper is the preferred way now to invoke exitcode from within a DAG.

exitcode This is the postscript that can be found at \$PEGASUS_HOME/bin/exitcode. It can be invoked on jobs that are launched via kickstart. This is a legacy postscript, and has now been replaced by exitpost.

none This means that no postscript is generated for the jobs. This is useful for MPI jobs that are not launched via kickstart currently.

any legal identifier Any other identifier of the form `([_A-Za-z][_A-Za-z0-9]*)`, than one of the 3 reserved keywords above, signifies a user postscript. This allows you to specify your own postscript for the jobs in the workflow. The path to the postscript can be specified by the property

```
pegasus.exitcode.path.[value]
```

where `[value]` is this legal identifier that you have specified. The arguments to be passed to your postscript can be set by the property `pegasus.exitcode.arguments`. If the path is not set for the user postscript, then the default path `$PEGASUS_HOME/bin/exitcode` is picked up.

The user postscript is passed the name of the `.out` file of the job as the last argument on the command line. For e.g. with the following property settings

```
pegasus.exitcode.default user_script
pegasus.exitcode.path.user_script /bin/user_postscript
pegasus.exitcode.arguments -verbose
```

for a job with submit file `a.sub` a user postscript invocation will appear in the `.dag` file as

```
/bin/user_postscript -verbose a.out
```

11.6 pegasus.exitcode.path.[value]

System	Pegasus
Type	string
Default	no default
Since	2.0
See also	pegasus.exitcode.impl, section 11.5 (page 36)
Old Name	vds.exitcode.path.[value]

This property allows you to override the default path to various postscripts. The `[value]` part of the property name needs to be replaced by any of the values of the `pegasus.exitcode.default` property.

For example, if you want to override the path to the user postscript and point to `/tmp/user_postscript`, you need to set the property `pegasus.exitcode.user.path` to `/tmp/user_postscript`. You can set paths to multiple postscripts. A single workflow of yours can use multiple postscripts for different types of jobs, by associating different Dagman profile key `POST` with the jobs.

For example you can have properties set as follows

```

pegasus.exitcode.path.exitcode = /my/development/pegasus/bin/exitcode
pegasus.exitcode.path.exitmode = /my/standard/pegasus/bin/exitpost
pegasus.exitcode.path.user1 = /a/b/c/d
pegasus.exitcode.path.user2 = /e/f/g/h

```

For some jobs you can choose to have exitcode invoked by associating Dagman profile key POST with value exitcode. For others, you can choose user1 invoked by associated Dagman profile key POST with value user1.

11.7 pegasus.exitcode.arguments

System	Pegasus
Type	string
Default	no default
Since	2.0
See also	pegasus.exitcode, section ?? (page ??)
Old Name	vds.exitcode.arguments

This specifies the arguments by which the exitcode is invoked on the kickstart output of a job. It applies to all the jobs, for which exitcode is invoked as determined by pegasus.exitcode property.

11.8 pegasus.exitcode.debug

System	Pegasus
Type	Boolean
Default	no default
Since	1.4.7
See also	pegasus.exitcode, section ?? (page ??)
Old Name	vds.exitcode.debug

This property turns on the debug for the exitcode and exitpost postscripts. By default the debug is turned off. On setting the property to true, an exitcode log file being created for each job in the submit directory. The log file created is of the format [JOBNAME].exit.log .

12 Remote Scheduler Properties

12.1 pegasus.remote.scheduler.projects

System	Remote Schedulers
Type	list of kv pairs
Default	(no default)
Example	jazz=PDQ,pnnl=emsl12491
Old Name	vds.scheduler.remote.projects

This property allows to set the **project** name that is to be used for each pool. Usually, such project names are specific to a small set of users, and can not be well set in the pool configuration file. Please note that the key is the pool handle, and the value is the project name. Setting the project will result in the generation of an RSL term (project=xxxx) for the matching pool handle.

The value specified by this property can be overridden by Globus profile key named project, that can be specified either in the transformation catalog, the site catalog or the DAX.

12.2 pegasus.remote.scheduler.queues

System	Remote Schedulers
Type	list of kv pairs
Default	(no default)
Example	jazz=PDQ,pnnl=emsl12491
Old Name	vds.scheduler.remote.queues

This property allows to set the **queue** name that is to be used for each pool. Usually, such queue names are specific to single users, and can thus not be well set in the pool configuration file. Please note that the key is the pool handle, and the value is the queue name. The property is applicable to any remote scheduling system that employs named queues, e.g. PBS or LSF. Setting the queue will result in the generation of an RSL clause (queue=xxxx) for the matching pool handle.

The value specified by this property can be overridden by Globus profile key named queue, that can be specified either in the transformation catalog, the site catalog or the DAX.

12.3 pegasus.remote.scheduler.min.[time]

System	Remote Schedulers
Type	list of kv pairs
Since	2.0
Default	(no default)
Example	jazz=10,pnnl=20
Old Name	vds.scheduler.remote.maxwalltimes, vds.scheduler.remote.min.[time]

This property allows to set a minimum max time for your jobs on each pool. This property is applicable to any remote scheduling system that employs walltimes like PBS. Setting the walltime will result in the generation of an RSL (Resource Specification Language) clause (`maxwalltime=xxxx`) for the matching pool handle. Please note that most if all scheduling systems that use this kill the job if the jobs running time exceed the advertised walltime in the job description.

The [time] part of the property name needs to be replaced by the any of the following time RSL keys supported by Globus GRAM.

maxwalltime Explicitly set the maximum walltime for a single execution of the executable. The units is in minutes.

maxtime The maximum walltime or cputime for a single execution of the executable. Walltime or cputime is selected by the GRAM scheduler being interfaced. The units is in minutes.

maxcputime Explicitly set the maximum cputime for a single execution of the executable. The units is in minutes.

More information on Globus RSL can be found at http://www.globus.org/toolkit/docs/2.4/gram/gram_rsl_parameters.html

The value specified by this property can be overridden by Globus profile key named `maxwalltime`, that can be specified either in the transformation catalog, the site catalog or the DAX.

13 Interface To Condor And Condor Dagman

The Condor DAGMan facility is usually activate using the `condor_submit_dag` command. However, many shapes of workflows have the ability to either overburden the submit host, or overflow remote gatekeeper hosts. While DAGMan provides throttles, unfortunately these can only be supplied on the command-line. Thus, PEGASUS provides a versatile wrapper to invoke DAGMan, called `pegasus-submit-dag`. It can be configured from the command-line, from user- and system properties, and by defaults.

13.1 `pegasus.condor.arguments.quote`

System	Condor
Type	Boolean
Default	true
Since	2.0
Old Name	<code>pegasus.condor.arguments.quote</code>

This property determines whether to apply the new Condor quoting rules for quoting the argument string. The new argument quoting rules appeared in Condor 6.7.xx series. We have verified it for 6.7.19 version. If you are using an old condor at the submit host, set this property to false.

13.2 pegasus.condor.release

System	Condor
Type	non-negative integer
Default	3
Since	2.0
See also	pegasus.condor.remove, section 13.3 (page 41)
See also	pegasus.condor.retry, section ?? (page ??)
Old Name	vds.scheduler.condor.release

This property determines the number of release attempts that are written into the submit file. Condor will hold jobs on certain kinds of failures. Many known failing conditions are transient, thus, if the job is automatically release after a short time, it usually progresses.

13.3 pegasus.condor.remove

System	Condor
Type	non-negative integer
Default	3
See also	pegasus.condor.release, section 13.2 (page 41)
See also	pegasus.condor.retry, section ?? (page ??)
Since	2.0
Old Name	vds.scheduler.condor.remove

This property determines the number of hold attempts that happen before Condor removes the job from the queue. This value is tied to the number of release attempts. Hence, Pegasus enforces $\text{release} > \text{remove} \geq 0$. A value of zero indicates, that the user does not want the job to be removed from the queue. In that case, no `periodic_remove` statement would be constructed. If the property is not set, the value of `pegasus.condor.release` is used. If both the properties are unset, then the default value of 3 is used.

13.4 pegasus.condor.output.stream

System	Condor
Type	Boolean
Value[0]	false
Value[1]	true
Default	true
Since	2.0
See also	pegasus.condor.error.stream, section 13.5 (page 42)
Old Name	vds.scheduler.condor.output.stream

By default, GASS streams back stdout continuously. In this default mode, it will require additional filedescriptors on the gatekeeper. Recent versions of Globus and Condor allow the content of stdout

to be streamed back after the job is done. While no content is visible during the execution of the job, it saves precious gatekeeper resources.

13.5 pegasus.condor.error.stream

System	Condor
Type	Boolean
Value[0]	false
Value[1]	true
Default	true
Since	2.0
See also	pegasus.condor.output.stream, section 13.4 (page 41)
Old Name	vds.scheduler.condor.error.stream

By default, GASS streams back stderr continuously. In this default mode, it will require additional filedescriptors on the gatekeeper. Recent versions of Globus and Condor allow the content of stderr to be streamed back after the job is done. While no content is visible during the execution of the job, it saves precious gatekeeper resources.

13.6 pegasus.dagman.retry

System	Condor DAGMan
Type	non-negative integer
Since	2.0
Default	3
See also	pegasus.condor.release, section 13.2 (page 41)
See also	pegasus.condor.remove, section 13.3 (page 41)
Old Name	vds.scheduler.condor.retry

This property determines the number of attempts DAGMAN makes to execute a job in case of failure. In case of deferred planning this can be used to do replanning. A failure during the execution of a partition can trigger the execution of Pegasus via a prescript, that may result in a different plan for the partition being generated and executed. Note, this is different from pegasus.scheduler.condor.release which results in condor retrying the same job (submit file) when a job goes in HOLD state.

13.7 pegasus.dagman.maxpre

System	DAGman wrapper
Type	integer
Default	20
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit_dag.html

The `pegasus-submit-dag` wrapper processes properties to set DAGMan commandline arguments. The argument sets the maximum number of PRE scripts within the DAG that may be running at one time. If this option is set to 0, the default number of PRE scripts is unlimited. The PEGASUS system throttles artificially to a maximum of 20 PRE scripts.

13.8 `pegasus.dagman.maxpost`

System	DAGman wrapper
Type	integer
Default	20
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit_dag.html

The `pegasus-submit-dag` wrapper processes properties to set DAGMan commandline arguments. The argument sets the maximum number of POST scripts within the DAG that may be running at one time. If this option is set to 0, the default number of POST scripts is unlimited. The PEGASUS system throttles artificially to a maximum of 20 POST scripts.

13.9 `pegasus.dagman.maxjobs`

System	DAGman wrapper
Type	integer
Default	200
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit_dag.html

The `pegasus-submit-dag` wrapper processes properties to set DAGMan commandline arguments. The argument sets the maximum number of jobs within the DAG that will be submitted to Condor at one time. If the option is omitted, the default number of jobs is unlimited. The PEGASUS system throttles artificially to a maximum of 200 simultaneous jobs that are visible to the Condor system. You may want to raise this bar.

13.10 `pegasus.dagman.nofity`

System	DAGman wrapper
Type	Case-insensitive enumeration
Value[0]	Complete
Value[1]	Error
Value[2]	Never
Default	Error
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit_dag.html
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit.html

The pegasus-submit-dag wrapper processes properties to set DAGMan commandline arguments. The argument sets the e-mail notification for DAGMan itself. This information will be used within the Condor submit description file for DAGMan. This file is produced by the the condor_submit_dag. See notification within the section of submit description file commands in the condor_submit manual page for specification of value. Many users prefer the value NEVER.

13.11 pegasus.dagman.verbose

System	DAGman wrapper
Type	Boolean
Value[0]	false
Value[1]	true
Default	false
Document	http://www.cs.wisc.edu/condor/manual/v6.6/condor_submit_dag.html

The pegasus-submit-dag wrapper processes properties to set DAGMan commandline arguments. If set and true, the argument activates verbose output in case of DAGMan errors.

14 Workflow Partitioning And Job Clustering Properties

14.1 pegasus.partitioner.label.key

System	Partitioner
Type	String
Default	label
Since	2.0
Old name	vds.label.key, vds.partitioner.label.key
See also	pegasus.clusterer.label.key, section 14.6 (page 47)

While partitioning the workflow into partitions, you can optionally label your graph to control how partitioning happens using a label based partitioning scheme. This is done by associating a profile key in the PEGASUS namespace with the jobs in the DAX. Each job that has the same value/label for this profile key, is put in the same partition.

This property allows you to specify the PEGASUS profile key that you want to use for label based partitioning.

14.2 pegasus.partitioner.horizontal.collapse.[txname]

System	Partitioner
Type	Integer
Default	no default
Since	2.0
See also	pegasus.partitioner.horizontal.bundle.[txname], section 14.3 (page 45)
Old Name	vds.partitioner.horizontal.bundle.[txname]

In case of horizontal partitioning, you can specify a "collapse" factor for transformations. The collapse factor for a transformation specifies the number of jobs referring to that transformation at a level that go into a single partition.

This property allows you to specify the collapse factors transformations where the [txname] in the property has to be replaced by the logical name of the transformation.

For example, if in a workflow you have 10 jobs referring to a transformation with a logical name txA at a particular level(l) of the workflow, and you have properties set as follows

```
pegasus.partitioner.horizontal.collapse.txA 4
```

The above would result 3 partitions being created, 2 having 4 jobs each and one having 2 jobs for that level (l).

14.3 pegasus.partitioner.horizontal.bundle.[txname]

System	Partitioner
Type	Integer
Default	no default
Since	2.0
See also	pegasus.partitioner.horizontal.collapse.[txname], section 14.2 (page 45)
Old Name	vds.partitioner.horizontal.bundle.[txname]

In case of horizontal partitioning, you can specify a "bundle" factor for transformations. The bundle factor for a transformation specifies the number of partitions that are created for each level of the workflow containing jobs that refer to that transformation.

This property allows you to specify the bundle factors transformations where the [txname] in the property has to be replaced by the logical name of the transformation.

For example, if in a workflow you have 10 jobs referring to a transformation with a logical name txA at a particular level(l) of the workflow, and you have properties set as follows

```
pegasus.partitioner.horizontal.bundle.txA 4
```

The above would result 4 partitions being created, 2 having 2 jobs each and 2 having 3 jobs for that level (l).

14.4 pegasus.clusterer.job.aggregator

System	Job Clustering
Since	2.0
Type	String
Value[0]	seqexec
Value[1]	mpiexec
Default	seqexec
Old Name	vds.job.aggregator

A large number of workflows executed through the Virtual Data System, are composed of several jobs that run for only a few seconds or so. The overhead of running any job on the grid is usually 60 seconds or more. Hence, it makes sense to collapse small independent jobs into a larger job. This property determines, the executable that will be used for running the larger job on the remote site.

seqexec In this mode, the executable used to run the merged job is seqexec that runs each of the smaller jobs sequentially on the same node. The executable "seqexec" is a PEGASUS tool distributed in the PEGASUS worker package, and can be usually found at {pegasus.home}/bin/seqexec.

mpiexec In this mode, the executable used to run the merged job is mpiexec that runs the smaller jobs via mpi on n nodes where n is the nodecount associated with the merged job. The executable "mpiexec" is a PEGASUS tool distributed in the PEGASUS worker package, and can be usually found at {pegasus.home}/bin/mpiexec.

14.5 pegasus.clusterer.job.aggregator.seqexec.hasgloballog

System	Job Clustering
Type	Boolean
Default	true
Since	2.0
See also	pegasus.clusterer.job.aggregator, section 14.4 (page 46)
Old Name	vds.job.aggregator.seqexec.isgloballog

Seqexec logs the progress of the jobs that are being run by it in a progress file on the remote cluster where it is executed. The progress log is useful for you to track the progress of your computations and remote grid debugging. The progress log file can be shared by multiple seqexec jobs that are running on a particular cluster as part of the same workflow. Or it can be per job.

This property sets the Boolean flag, that indicates whether to have a single global log for all the seqexec jobs on a particular cluster or progress log per job.

14.6 pegasus.clusterer.label.key

System	Job Clustering
Type	String
Default	label
Since	2.0
See also	pegasus.partitioner.label.key, section 14.1 (page 44)
Old Name	vds.clusterer.label.key

While clustering jobs in the workflow into larger jobs, you can optionally label your graph to control which jobs are clustered and to which clustered job they belong. This is done using a label based clustering scheme and is done by associating a profile/label key in the PEGASUS namespace with the jobs in the DAX. Each job that has the same value/label value for this profile key, is put in the same clustered job.

This property allows you to specify the PEGASUS profile key that you want to use for label based clustering.

15 Miscellaneous Properties

15.1 pegasus.job.priority

System	Pegasus
Type	Integer
Default	no default
Since	2.0
See also	pegasus.transfer.*.priority, section 10.14 (page 33)
Old Name	vds.job.priority

This property sets the priority for all the condor jobs. The priority appears as a classad in the condor submit files. Optionally, user can also set the priorities for the transfer jobs separately by specifying the property pegasus.transfer.*.priority.

15.2 pegasus.catalog.transformation.mapper

System	Staging of Executables
Since	2.0
Type	enumeration
Value[0]	All
Value[1]	Installed
Value[2]	Staged
Value[3]	Submit
Default	All
See also	pegasus.transformation.selector, section ?? (page ??)
Old Name	vds.tc.mapper

Pegasus now supports transfer of statically linked executables as part of the concrete workflow. At present, there is only support for staging of executables referred to by the compute jobs specified in the DAX file. Pegasus determines the source locations of the binaries from the transformation catalog, where it searches for entries of type `STATIC_BINARY` for a particular architecture type. The PFN for these entries should refer to a globus-url-copy valid and accessible remote URL. For transfer of executables, Pegasus constructs a soft state map that resides on top of the transformation catalog, that helps in determining the locations from where an executable can be staged to the remote site.

This property determines, how that map is created.

All In this mode, all sources with entries of type `STATIC_BINARY` for a particular transformation are considered valid sources for the transfer of executables. This the most general mode, and results in the constructing the map as a result of the cartesian product of the matches.

Installed In this mode, only entries that are of type `INSTALLED` are used while constructing the soft state map. This results in Pegasus never doing any transfer of executables as part of the workflow. It always prefers the installed executables at the remote sites.

Staged In this mode, only entries that are of type `STATIC_BINARY` are used while constructing the soft state map. This results in the concrete workflow referring only to the staged executables, irrespective of the fact that the executables are already installed at the remote end.

Submit In this mode, only entries that are of type `STATIC_BINARY` and reside at the submit host (pool local), are used while constructing the soft state map. This is especially helpful, when the user wants to use the latest compute code for his computations on the grid and that relies on his submit host.

15.3 pegasus.selector.transformation

System	Staging of Executables
Since	2.0
Type	enumeration
Value[0]	Random
Value[1]	Installed
Value[2]	Staged
Value[3]	Submit
Default	Random
See also	pegasus.catalog.transformation, section 7.3.1 (page 19)
Old Name	vds.transformation.selector

In case of transfer of executables, Pegasus could have various transformations to select from when it schedules to run a particular compute job at a remote site. For e.g it can have the choice of staging an executable from a particular remote pool, from the local (submit host) only, use the one that is installed on the remote site only.

This property determines, how a transformation amongst the various candidate transformations is selected, and is applied after the property pegasus.tc has been applied. For e.g specifying pegasus.tc as Staged and then pegasus.transformation.selector as INSTALLED does not work, as by the time this property is applied, the soft state map only has entries of type STAGED.

Random In this mode, a random matching candidate transformation is selected to be staged to the remote execution pool.

Installed In this mode, only entries that are of type INSTALLED are selected. This means that the concrete workflow only refers to the transformations already pre installed on the remote pools.

Staged In this mode, only entries that are of type STATIC_BINARY are selected, ignoring the ones that are installed at the remote site.

Submit In this mode, only entries that are of type STATIC_BINARY and reside at the submit host (pool local), are selected as sources for staging the executables to the remote execution pools.

15.4 pegasus.auth.gridftp.timeout

System	Pegasus
Since	2.0
Type	non negative integer
Default	120
Old Name	vds.auth.gridftp.timeout

This property sets the socket timeout on the socket that is opened to the remote gridftp server for authentication purposes.

Index

pegasus.auth.gridftp.timeout, 49
pegasus.catalog.*.db.*, 12
pegasus.catalog.*.db.password, 12
pegasus.catalog.*.db.url, 11
pegasus.catalog.*.db.user, 12
pegasus.catalog.<catalog-name>.db.driver, 11
pegasus.catalog.provenance, 14
pegasus.catalog.replica, 15
pegasus.catalog.replica.cache.asrc, 18
pegasus.catalog.replica.chunk.size, 17
pegasus.catalog.replica.lrc.ignore, 17
pegasus.catalog.replica.lrc.restrict, 17
pegasus.catalog.replica.url, 16
pegasus.catalog.site, 18
pegasus.catalog.site.file, 19
pegasus.catalog.transformation, 19
pegasus.catalog.transformation.file, 20
pegasus.catalog.transformation.mapper, 48
pegasus.catalog.work, 15
pegasus.clusterer.job.aggregator, 46
pegasus.clusterer.job.aggregator.seqexec.hasgloballog, 46
pegasus.clusterer.label.key, 47
pegasus.condor.arguments.quote, 40
pegasus.condor.error.stream, 42
pegasus.condor.output.stream, 41
pegasus.condor.release, 41
pegasus.condor.remove, 41
pegasus.dagman.maxjobs, 43
pegasus.dagman.maxpost, 43
pegasus.dagman.maxpre, 42
pegasus.dagman.notify, 43
pegasus.dagman.retry, 42
pegasus.dagman.verbose, 44
pegasus.dir.create, 9
pegasus.dir.exec, 8
pegasus.dir.storage, 8
pegasus.exitcode.arguments, 38
pegasus.exitcode.debug, 38
pegasus.exitcode.impl, 36
pegasus.exitcode.path.[value], 37
pegasus.exitcode.scope, 35
pegasus.gridstart, 34
pegasus.gridstart.invoke.always, 34
pegasus.gridstart.invoke.length, 35
pegasus.home.datadir, 7
pegasus.home.localstatedir, 8
pegasus.home.sharedstatedir, 7
pegasus.home.sysconfdir, 7
pegasus.job.priority, 47
pegasus.partitionner.horizional.bundle.[txname], 45
pegasus.partitionner.horizional.collapse.[txname], 45
pegasus.partitionner.label.key, 44
pegasus.properties, 6
pegasus.remote.scheduler.min.[time], 39
pegasus.remote.scheduler.projects, 39
pegasus.remote.scheduler.queues, 39
pegasus.schema.dax, 9
pegasus.schema.ivr, 10
pegasus.schema.pdax, 10
pegasus.schema.sc, 10
pegasus.selector.replica, 21
pegasus.selector.replica.*.ignore.stagein.sites, 22
pegasus.selector.replica.*.prefer.stagein.sites, 22
pegasus.selector.site, 23
pegasus.selector.site.keep.tmp, 25
pegasus.selector.site.path, 24
pegasus.selector.site.timeout, 25
pegasus.selector.transformation, 49
pegasus.site.selector.env.*, 25
pegasus.transfer.*.impl, 26
pegasus.transfer.*.priority, 33
pegasus.transfer.*.thirdparty.remote, 32
pegasus.transfer.*.thirdparty.sites, 31
pegasus.transfer.arguments, 29
pegasus.transfer.disable.chmod.sites, 33
pegasus.transfer.force, 30

pegasus.transfer.links, 29
pegasus.transfer.proxy, 33
pegasus.transfer.refiner, 28
pegasus.transfer.single.quote, 30
pegasus.transfer.staging.delimiter, 32
pegasus.transfer.throttle.processes, 30
pegasus.transfer.throttle.streams, 31
pegasus.user.properties, 6